

Evaluation of search-enabled Pre-trained Large Language Models on retrieval tasks for the PubChem Database

Ash Sze and Soha Hassoun
Tufts University

October 18, 2024

Supporting File 1

Detailed Results on the Protocol

To provide additional details for each protocol, we describe our experiences with GPT-4o in more detail, highlighting values and limitations that are specific to each case. Please note that Protocol #3 and #4 are detailed as examples in the main manuscript.

Protocol 1. The first protocol identifies genes and proteins that interact with losartan. Such data in PubChem originates from several sources including DrugBank (?), Drug-Gene Interaction database (?), and others. The PubChem web interface guides the user to go to section 16.2 (Chemical-Target Interaction), and filter against DrugBank interactions. The gold results (see full results in appendix) include several Cytochrome enzymes, UDP-glucuronosyltransferases, and others.

Based on prompting, GPT-4o provides correct yet incomplete results. The response is missing certain entries, e.g., Cytochrome P450 2C19, Cytochrome P450 2C8, and ATO- dependent translocase ABCB1. GPT-4o elaborates on each interaction with information external to PubChem and DrugBank such as Wikipedia (see partial results in Figure S1). The names of some interactions and proteins differ slightly from PubChem, indicating the results were acquired from external sources or training data. While the PubChem protocol had a strict definition of what qualifies as an interaction, and where the interactions must be sourced from, GPT-4o interpreted interactions as broader, and therefore included information about downstream effects of these interactions, listing “RAAS” and “Plasma Renin Activity” in the list of proteins and genes affected. Enhanced prompting for this protocol involved excluding downstream effects and protein or gene information and emphasizing the need for data to only be from PubChem. Unfortunately, any engineering to enhance the prompt

led to incorrect answers with non-PubChem sources which strayed further from the established gold answer.

There is currently no method to retrieve the same information through programmatic access. Therefore, the silver answer is not available.

Protocol 2. The second protocol aims to find drug-like compounds similar to losartan based on a two-dimensional (2-D) similarity search using PubChem. The PubChem web interface guides users to perform a "Similar Structures Search" and apply Lipinski's rule of 5 filters. The gold answer includes a detailed list of compounds that match the specified criteria (see full results in the appendix).

GPT-4o attempts to provide a list of compounds that are structurally similar to losartan while adhering to Lipinski's rule of five. However, the results are incorrect when compared to the gold answer. None of GPT-4o's outputs, such as Candesartan and Irbesartan, match the outputs from the gold answer. Looking into the external links provided by GPT-4o show that the model interpreted losartan as all types of losartan, including losartan Sodium. Enhanced prompting for this protocol involved specifying losartan as only "losartan" and not other variations of the drug. This enhanced prompt only shortened the output, leaving it with Candesartan, Irbesartan, and Valsartan, which were still incorrect compared to the gold standard.

To retrieve similar information programmatically, we used PUG to perform a 2-D substructure search for losartan with CID 3961 and applied filters for Lipinski's rule of 5. The programmatic access involved using a URL to perform the search and obtain a list key, which was then used to request a detailed list of hit CIDs. This method provided comprehensive and accurate results that align with PubChem's standards, but upon comparison with the gold answer, the retrieved information filtered on broader criteria, leading to a longer list of hit targets inclusive of the gold answer. This silver answer highlights a significant difference between searches done on PUG and the PubChem web interface.

We also prompted GPT-4o with a programmatic access approach to provide a PUG link for the 2-D similarity search. The prompt included specific criteria, such as molecular weight, hydrogen bond donors and acceptors, and partition coefficient limits according to Lipinski's rule of five. GPT-4o generated the silver answer PUG URL. Although this is a multi-step protocol, requiring the user to manually input the generated list key for the final result, the programmatic prompt effectively provided the most relevant PUG URL templates for each step for detailed data access (refer to Figure S2 for a full protocol summary).

Protocol 5. The fifth protocol aims to find drugs that target the human type-1 angiotensin II receptor gene. The data originates from several sources within PubChem, including DrugBank. The PubChem web interface guides users to search for "type 1 angiotensin II receptor," select the "Genes" tab and click the entry with "Human," navigate to the 7.1 "Chemical-Gene Interactions" subsection, sort by "Data Source" to view all interacting compounds from DrugBank, and download the necessary sections (see full results in the appendix).

GPT-4o provides a list of drugs that interact with the gene encoding the

human type-1 angiotensin II receptor. The response is correct but incomplete. GPT-4o lists the drugs Valsartan, Olmesartan, Telmisartan, Irbesartan, Candesartan, and Azilsartan medoxomil, which is only 7 out of the 14 total drug interactions from DrugBank. Attempting to enhance the prompt by asking for a complete list, or all of the data lead to faulty code generation or an output unchanged from the initial output.

There is no programmatic way to directly access the gene-chemical interactions table. There is no silver response.

Protocol 6. The sixth protocol aims to get the bioactivity data of all chemicals tested against the human type-1 angiotensin II receptor (AT1R) and its rat ortholog. The data originates from various sources within PubChem. The PubChem web interface guides users to access the protein summary page for “type 1 angiotensin II receptor,” navigate to the “tested compounds” section, download the list of tested compounds and their bioactivity data, and repeat the process for “P29089 (Norway rat)” under the orthologous protein section (see full results in the appendix).

GPT-4o provides a summary of the bioactivity data for the most important chemicals tested against the human type-1 angiotensin II receptor and its rat ortholog. GPT-4o cannot directly search PubChem and thus only provides possible bioactivity types, a general overview of the AT1R receptor, and listings of human AT1R interactions for compounds losartan, Olmesartan, and Candesartan, and rat AT1R interactions for compounds Valsartan and Irbesartan. GPT-4o does not explain how it shortlisted these compounds, and descriptions of the compounds are generic and do not reflect the gold answer. Enhanced prompting by re-querying and asking for exact lists of compounds did not improve the quality of the responses for this protocol.

To retrieve similar information programmatically, we utilized PUG to perform a concise bioactivities search using the gene IDs 185 and 81638, corresponding to AGTR1 and AGTR1b respectively. This method involved finding the relevant gene IDs and using URLs to obtain detailed bioactivity data.

We also prompted GPT-4o with a programmatic access approach to provide PUG links for concise gene searches on gene ID 185 and gene ID 81638. The prompt included specific criteria to return a summary of bioactivity data for these gene IDs. GPT-4o generated incorrect PUG URLs due to an invalid input domain (refer to figure S4 for a full protocol summary).

Protocol 7. The seventh protocol aims to find compounds annotated with specific classifications or ontological terms using PubChem. The data originates from PubChem’s chemical database. The PubChem web interface guides users to the Classification Browser, selecting “MeSH” for classification and “Compound” for data, searching for “Antihypertensive agents” and saving as “1,” repeating for “anti-arrhythmia agents” and saving as “2,” and using the AND operator to search for results matching both classifications (see full results in the appendix).

GPT-4o cannot perform a direct search on PubChem and is thus unable to use the Boolean search function required for this protocol. The original protocol requires the user to create two separate searches and use the Boolean

operator AND to find results that match both statements. While PubChem has 68 results for this search, none of them correspond to GPT-4o’s answers. GPT-4o instead generates hypothetical datasets of chemicals and interactions, which do not accurately reflect the actual data available in PubChem. Enhanced prompting did not improve the quality of the responses for this protocol.

There is currently no way to programmatically access classifications or search by MeSH annotations via PUG-REST. The PUG-REST tutorial acknowledges this as a feature that may be added in the future based on demand. Thus, there is no silver answer for this protocol through programmatic access.

Protocol 8. The eighth protocol aims to find stereoisomers and isotopomers of a compound using an identity search in PubChem. The data originates from PubChem’s comprehensive chemical database. The PubChem web interface guides users to search for “CID 60846 structure,” select “Identity” and select “Same Isotope” in “Settings” for stereoisomers and “Same stereo” for stereoisotopomers and download the results as a CSV file (see full results in the appendix).

GPT-4o provides a list of stereoisomers and isotopomers for Valsartan (CID 60846). The response includes the (S)-enantiomer as a commonly known stereoisomer, and Valsartan-d3 and [3H]Valsartan as isotopomers. However, the names and formatting of these chemical compounds differ from PubChem’s data, indicating discrepancies. The response does not match the gold answer from PubChem. Attempting to enhance the prompt by changing “find” to “list” or “retrieve” led to faulty code generation. Re-querying for the correct answer leads to a change in output compounds that was still incorrect.

To retrieve similar information programmatically, we used PUG to perform a fast identity structure search for stereoisomers and isotopomers using the CID 60846. This method involved setting “identity_type” as “same_stereo” for stereoisotopomers and “same_isotope” for stereoisomers and using specific URLs to obtain the silver answer which matched the gold answer.

We prompted GPT-4o with a programmatic access approach to provide PUG links for a fast identity search using the same stereo and same isotope settings for CID 60846. However, the output generated by GPT-4o was incorrect, resulting in PUG-REST faults due to invalid operations. (refer to figure S6 for a full protocol summary)

<p>Reference Protocol</p> <ol style="list-style-type: none"> 1. Visit the PubChem homepage 2. Search the chemical name "losartan" 3. View the summary page of the best match 4. Go to section 16.2, "Chemical-Target Interactions" and filter for "DrugBank" interactions 5. Click download to view the full of list of macromolecules interacting with losartan 	<p>Programmatic Access Protocol</p> <p>There is currently no method to programmatically retrieve the "Chemical-Target Interactions" table data directly via PUG. There is no silver answer for this protocol.</p>
<p>GPT Generation via Search</p> <p>GPT Prompt: "Based only on information from PubChem, identify genes and proteins that interact with the compound losartan. Please provide a list of genes and proteins known to interact with losartan, including their roles in the interaction"</p> <p>GPT output: List including Angiotensin II Type 1 Receptor, Cytochrome P450 Enzymes, UGT Enzymes, Renin-Angiotensin-Aldosterone System, Plasma Renin Activity.</p>	<p>GPT Generation via Programmatic Prompt</p> <p>GPT Programmatic Prompt and PUG-based GPT output is not applicable because there is no programmatic access method or silver answer.</p>

Figure 1: Protocol #1 for finding genes and proteins that interact with a given compound, made specific for losartan.

<p>Reference Protocol</p> <ol style="list-style-type: none"> 1. Visit the PubChem homepage 2. Search the chemical name "losartan" 3. Click for "Similar Structures Search" 4. Through "Settings", set "Filters" for Lipinski's rule of 5 5. Click download to view the full hit list 	<p>Programmatic Access Protocol</p> <p>There is a programmatic way to perform a 2D structure search with limited filters compared to the PubChem web interface. The silver answer is the broader 2D structure search results.</p> <ol style="list-style-type: none"> 1. Use PUG to perform a 2D substructure search with losartan CID 3961 and filter by Lipinski's rules via URL operations. The return value is a list key. https://pubchem.ncbi.nlm.nih.gov/rest/pug/compound/similarity/cid/3961/JSON?MaxMW=500&HBD=5&HBA=10&LogP=5 2. Use the generated list key to perform a request for a list of hit CIDs. https://pubchem.ncbi.nlm.nih.gov/rest/pug/compound/listkey/{listkey}/cids/JSON
<p>GPT Generation via Search</p> <p>GPT Prompt: "Based only on information from Pubchem, find drug-like compounds that are structurally similar to the compound losartan based on two-dimensional (2-D) similarity that satisfy Lipinski's rule of five. (Lipinski's rule of 5 listed and specified)"</p> <p>GPT output: List including Candesartan, Irbesartan, Valsartan, Eprosartan, Telmisartan.</p>	<p>GPT Generation via Programmatic Prompt</p> <p>GPT Programmatic Prompt: "Provide the PUG URL for a 2D similarity search on CID 3961 with the following filters: molecular weight less than 500 g/mol, No more than 5 hydrogen bond donors, No more than 10 hydrogen bond acceptors, An octanol-water partition coefficient (log P) that does not exceed 5. Provide the URL template to load the list key."</p> <p>PUG-based GPT output: Same as silver answer.</p>

Figure 2: Protocol #2 for finding drug-like compounds similar to a query compound through 2-D similarity search, made specific for losartan filtered through Lipinski's rule of 5.

<p>Reference Protocol</p> <ol style="list-style-type: none"> 1. Visit the PubChem homepage 2. Search for "type 1 angiotensin II receptor" 3. Select the "Genes" tab and click the entry with "Human" 4. Go to the 7.1 "Chemical-Gene Interactions" subsection 5. Sort by "Data Source" to view all interacting compounds from DrugBank and download sections as needed. 	<p>Programmatic Access Protocol</p> <p>There is currently no method to programmatically retrieve the "Chemical-Gene Interactions" table data directly via PUG. There is no silver answer for this protocol.</p>
<p>GPT Generation via Search</p> <p>GPT Prompt: "Based on information only from PubChem, find the most important drugs that interact with the gene encoding the human type-1 angiotensin II receptor, which is the target of losartan."</p> <p>GPT output: List of drugs including Valsartan, Olmesartan, Telmisartan, Irbesartan, Candesartan, Eprosartan, and Azilsartan medoxomil.</p>	<p>GPT Generation via Programmatic Prompt</p> <p>GPT Programmatic Prompt and PUG-based GPT output is not applicable because there is no programmatic access method or silver answer.</p>

Figure 3: Protocol #5 for finding drugs that target a particular gene, made specific for the gene type 1 angiotensin receptor (human).

<p>Reference Protocol</p> <ol style="list-style-type: none"> 1. Access the protein summary page for "type 1 angiotensin II receptor" similar to protocol 5 2. Go to the "tested compounds" section 3. Download the list of tested compounds and their bioactivity data 4. Repeat the process for "P29089(Norway rat)" under the orthologous protein section. 	<p>Programmatic Access Protocol</p> <p>There is a programmatic way to access the bioactivity data of chemicals tested against a protein. The silver answer is more detailed than the table found through PubChem</p> <ol style="list-style-type: none"> 1. Find the gene id for AGTR1 and AGTR1b (185 and 81638) 2. Perform a concise bioactivities search using the gene IDs <p>https://pubchem.ncbi.nlm.nih.gov/rest/pug/gene/geneid/185/concise/JSON https://pubchem.ncbi.nlm.nih.gov/rest/pug/gene/geneid/81638/concise/JSON</p>
<p>GPT Generation via Search</p> <p>GPT Prompt: "Based on information only from PubChem, create or link to a summary for and list the bioactivity data of the most important chemicals tested against the human type-1 angiotensin II receptor and its rat ortholog."</p> <p>GPT output: List of summarized bioactivity data for human AT1R including losartan, Olmesartan, and Cadesartan, and for rat AT1R including Valsartan, and Irbesartan.</p>	<p>GPT Generation via Programmatic Prompt</p> <p>GPT Programmatic Prompt: "Provide the PUG URL for a concise gene search on gene ID 185 and another URL for gene ID 81638"</p> <p>Pug-based GPT Output: https://pubchem.ncbi.nlm.nih.gov/rest/pug/genes/geneid/185/summary/JSON (PUG-REST fault due to invalid input domain) https://pubchem.ncbi.nlm.nih.gov/rest/pug/genes/geneid/81638/summary/JSON (PUGREST fault due to invalid input domain)</p>

Figure 4: Protocol #6 for getting bioactivity data of all chemicals tested against a protein, made specific for the protein type 1 angiotensin II receptor (human and Norway rat).

<p>Reference Protocol</p> <ol style="list-style-type: none"> 1. Visit the PubChem Classification Browser 2. Select "MeSH" for classification and "Compound" for data 3. Search for "Antihypertensive agents" and save as "1" 4. Repeat for "anti-arrhythmia agents" and save as "2" 5. Use the AND operator to search "1" and "2" then view results 	<p>Programmatic Access Protocol</p> <p>There is no way to programmatically access classifications or search by MeSH annotations via PUG-REST yet. The PUG-REST tutorial acknowledges this as a feature that may or more not be added in the future based on demand. There is no silver answer for this protocol.</p>
<p>GPT Generation via Search</p> <p>GPT Prompt: "Based only on information from PubChem, list the chemicals with the same therapeutic uses as losartan, based on the MeSH annotations"</p> <p>GPT output: Unable to directly perform the Boolean operations required to access the information, GPT is limited to creating hypothetical datasets for chemicals and interactions.</p>	<p>GPT Generation via Programmatic Prompt</p> <p>GPT Programmatic Prompt and PUG-based GPT output is not applicable because there is no programmatic access method or silver answer.</p>

Figure 5: Protocol #7 for finding compounds annotated with classifications or ontological terms, made specific for "antihypertensive agents" and "anti-arrhythmia agents."

<p>Reference Protocol</p> <ol style="list-style-type: none"> 1. Visit the PubChem homepage 2. Search for "CID 60846 structure" 3. Click "Identity" and select "Same Isotope" in "Settings" for stereoisomers and "Same stereo" for isotopomers 4. Download for a full CSV 	<p>Programmatic Access Protocol</p> <p>There is a way to directly access the stereoisomers and isotopomers of a given CID. The silver answer is the gold answer.</p> <ol style="list-style-type: none"> 1. Find the CID of the search compound (60846) 2. Perform a fast identity structure search using PUG, setting "identity_type" as "same_stereo" for stereo-isotopomers and "same_isotope" for stereoisomers. <p>https://pubchem.ncbi.nlm.nih.gov/rest/pug/compound/fastidentity/cid/60846/cids/TXT?identity_type=same_stereo https://pubchem.ncbi.nlm.nih.gov/rest/pug/compound/fastidentity/cid/60846/cids/TXT?identity_type=same_isotope</p>
<p>GPT Generation via Search</p> <p>GPT Prompt: "Based only on information from PubChem, find stereoisomers and isotopomers of a given compound, with valsartan (CID 60846)"</p> <p>GPT output: Listed (S)-enantiomer as a commonly known stereoisomer, and Valsartan-d3 and [3H]valsartan as isotopomers.</p>	<p>GPT Generation via Programmatic Prompt</p> <p>GPT Programmatic Prompt: "Provide the PUG URLs for a fast identity search for the same stereo and the same isotope on CID 60846"</p> <p>GPT Output: https://pubchem.ncbi.nlm.nih.gov/rest/pug/compound/fastidentity/cid/60846/same_stereo/JSON (PUG-REST fault due to Invalid Operation) https://pubchem.ncbi.nlm.nih.gov/rest/pug/compound/fastidentity/cid/60846/same_isotope/JSON (PUG-REST fault due to Invalid Operation)</p>

Figure 6: Protocol #8 for getting stereoisomers and isotopomers of a compound through identity search, made specific for a search on "CID 60846 structure."